

Using Machine Learning Tools

Deep Neural Networks

University of Adelaide

Previously ...

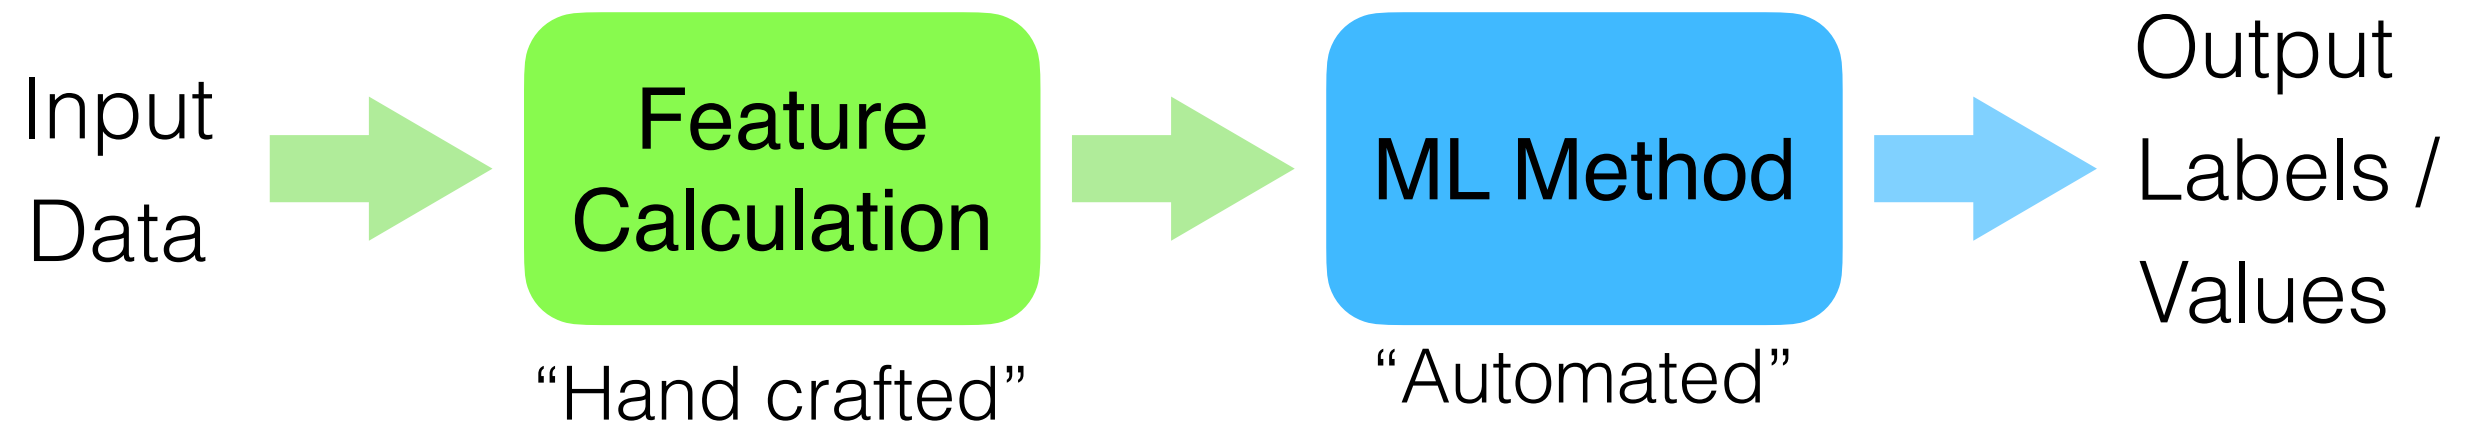
- Last time we discussed clustering, data visualisation and dimensionality reduction methods
- Prior to that we considered regression and classification ML
 - Defined loss (or cost) functions (e.g. MSE)
 - Discussed gradient-based optimisation methods (e.g. SGD)
 - Importance of learning rate and initialisation

Today

- Introduction to Neural Networks
- Network architecture: layers, connectivity and features
- Neurons: nonlinearities and activation functions
- Parameters: connections, weights and biases
- Number of parameters
- Loss functions, epochs, batches, optimisers and training
- Classification and regression variants


Introduction to Neural Networks

Traditional Machine Learning




Introduction to Neural Networks

Classical Machine Learning




Deep Learning




Neurons

Interconnected
Neurons



Single Neuron



Neurons

Inputs

x_1

w_1

x_2

w_2

x_3

w_3

\vdots

w_n

x_n

Score function


$$f = w \cdot x + b$$

$f(\dots)$

b

+1


Bias term



Activation
function

$\sigma(f)$

Output
 y



$$y = \sigma(w \cdot x + b)$$
$$= \sigma(\sum_j w_j \cdot x_j + b)$$

Neural Network


Input
Layer

Hidden
Layer 1

Hidden
Layer 2

Output
Layer

Fully
Connected
Layout




Neural Network

Input Layer Hidden Layer 1 Hidden Layer 2 Output Layer

Input values as 1D vector

Fully Connected Layout

Each layer *only* connects to next layer




Bias inputs are not shown but implicit

Neural Network

Input Layer Hidden Layer 1 Hidden Layer 2 Output Layer

Parameters
= # weights +
bias



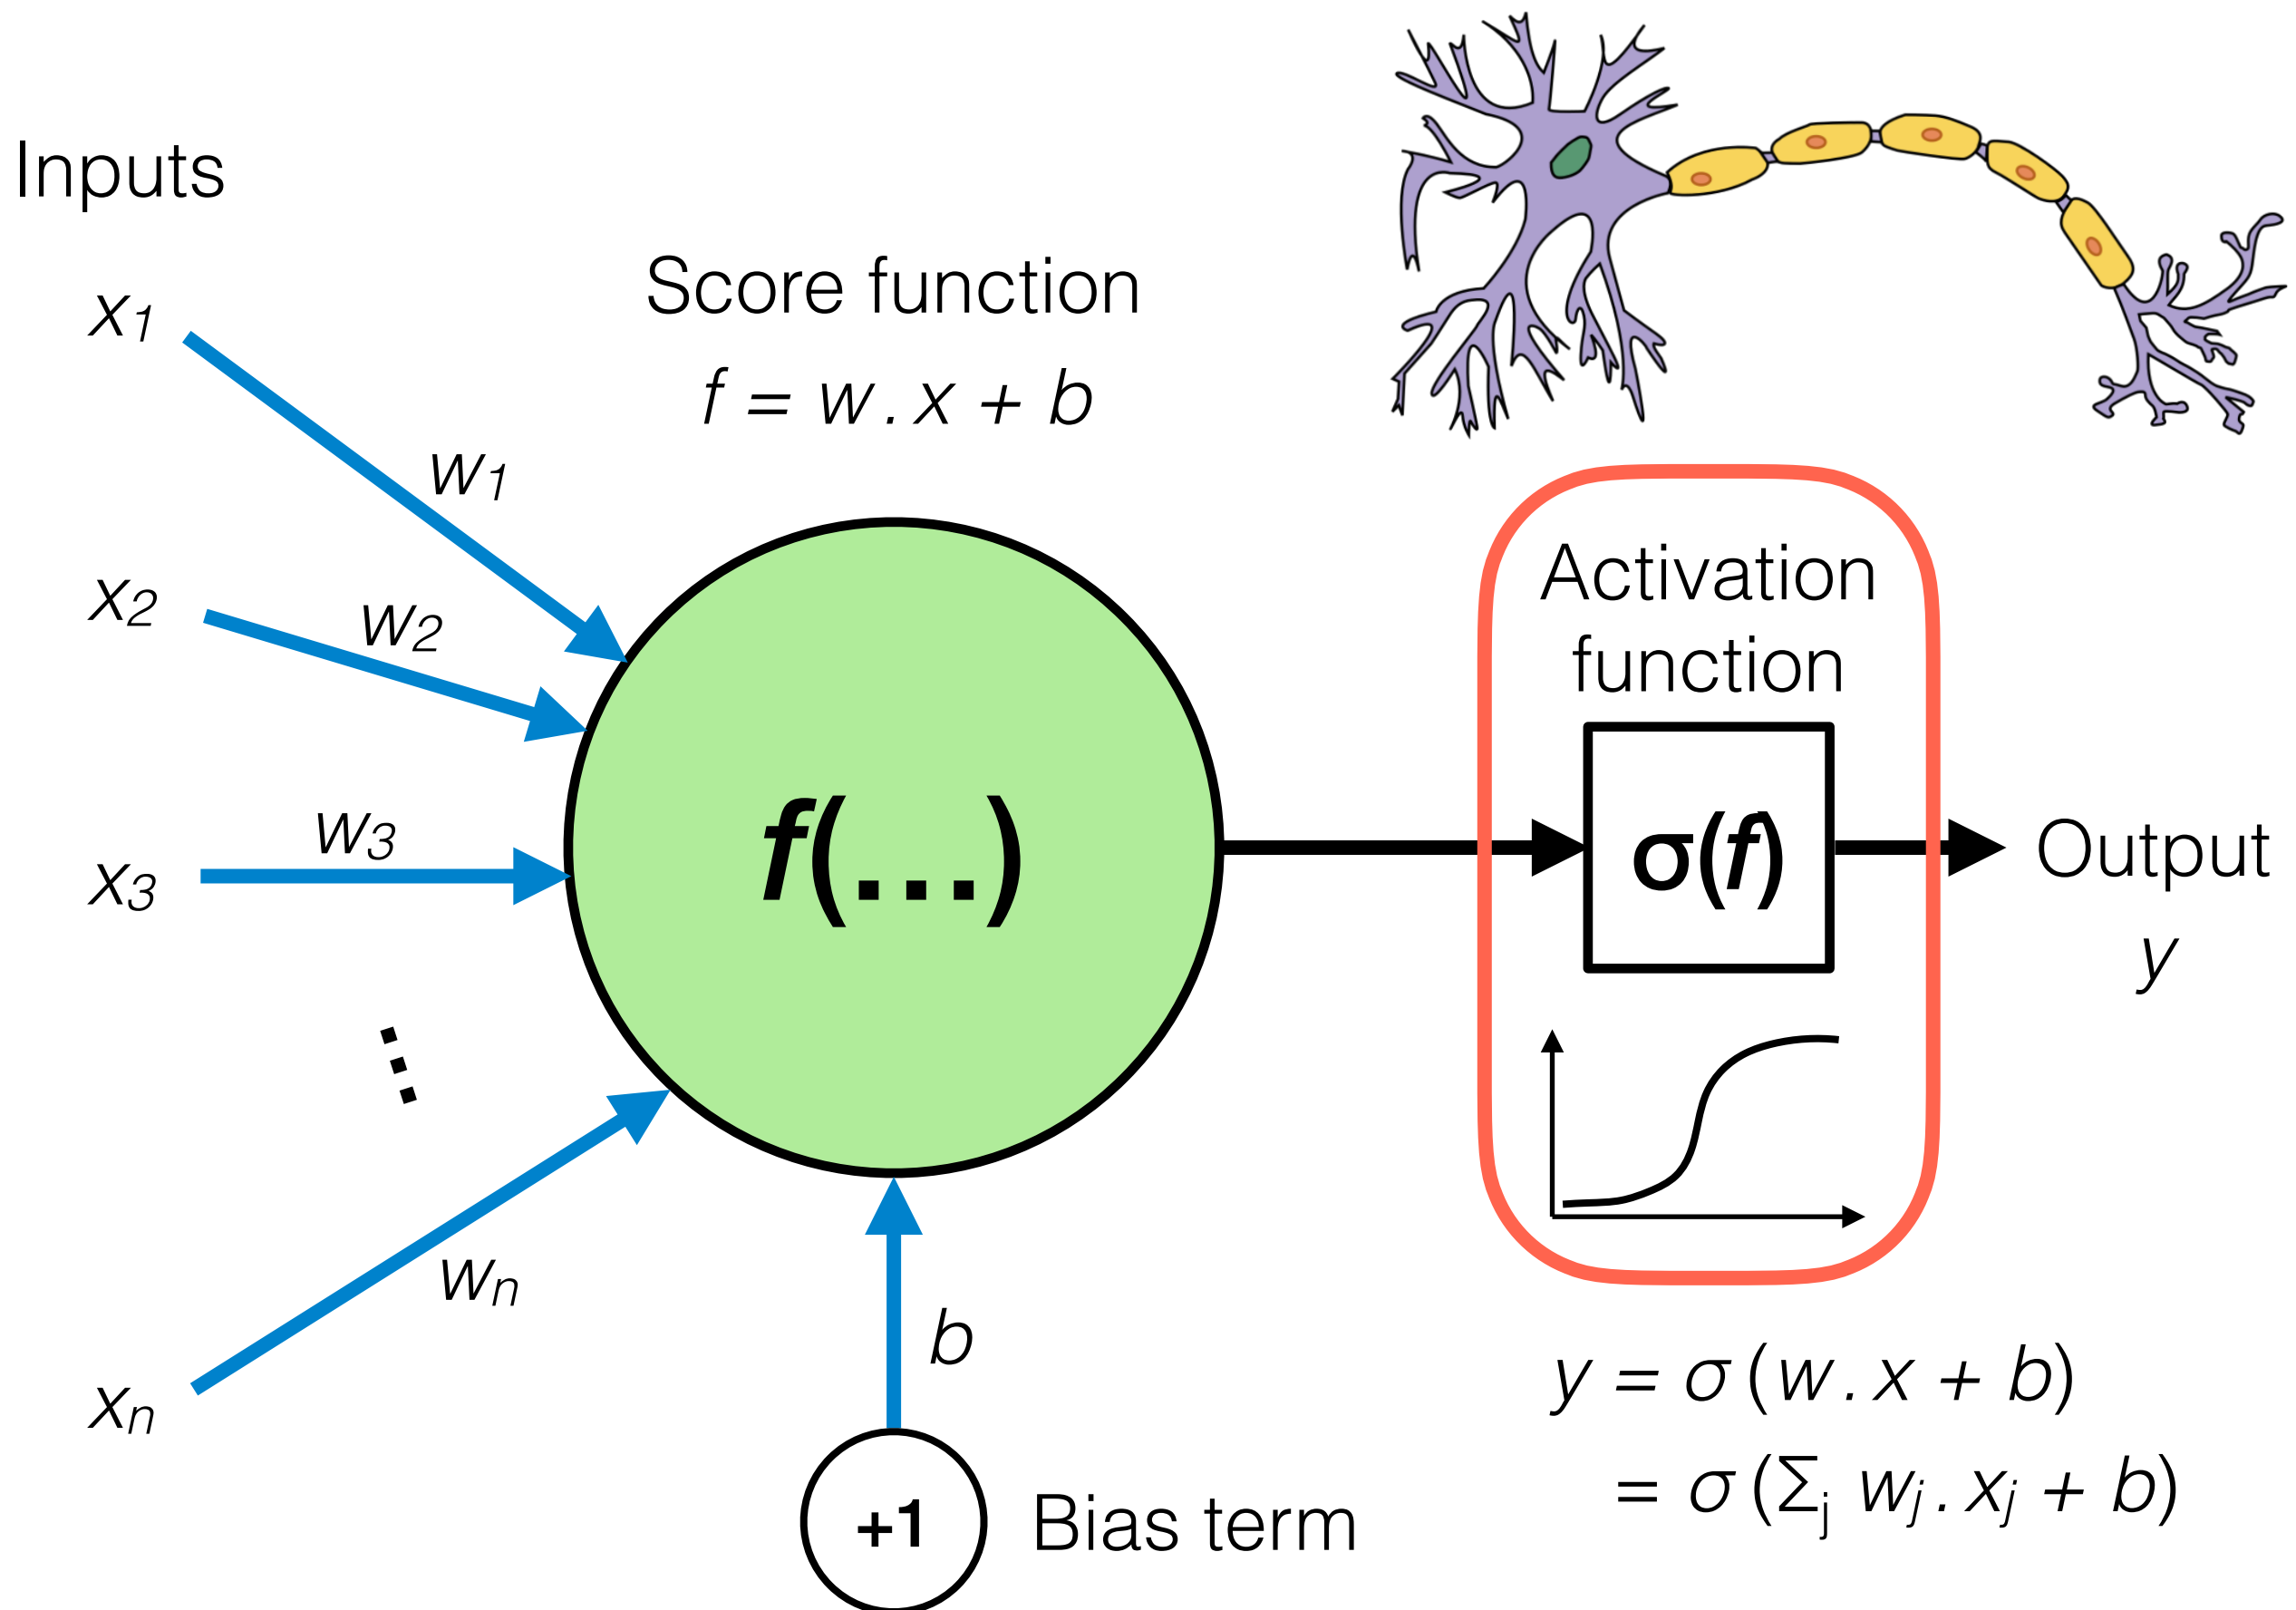
$$5 * 5 + 5 * 6 + 1 * 6 \\ = 61 \text{ parameters}$$

4 inputs +
1 bias

5 inputs +
1 bias




5 inputs +
1 bias

Importance of Nonlinearity



Importance of Nonlinearity


Without the nonlinearity the whole network just computes a linearly weighted sum and these are equivalent.



The nonlinearity is crucial, but a number of options exist, all similar to sigmoids.

ReLU (Rectified Linear Unit) is the most popular.

Loss Function



$$y = g(x; w, b)$$

$$\text{loss}(y_{\text{pred}}, y_{\text{train}})$$

$$y_{\text{pred}} = g(x_{\text{train}})$$

Minimise loss by
changing parameters
(weights and biases)

Sum loss over
batches/epochs

Optimisation


Optimisation method
used to find minimum
loss value

Typically use methods
based on gradient
descent:
backpropagation

Important to
have right:

- initialisation
- learning rate


Loss function



Optimisation: Batches and Epochs

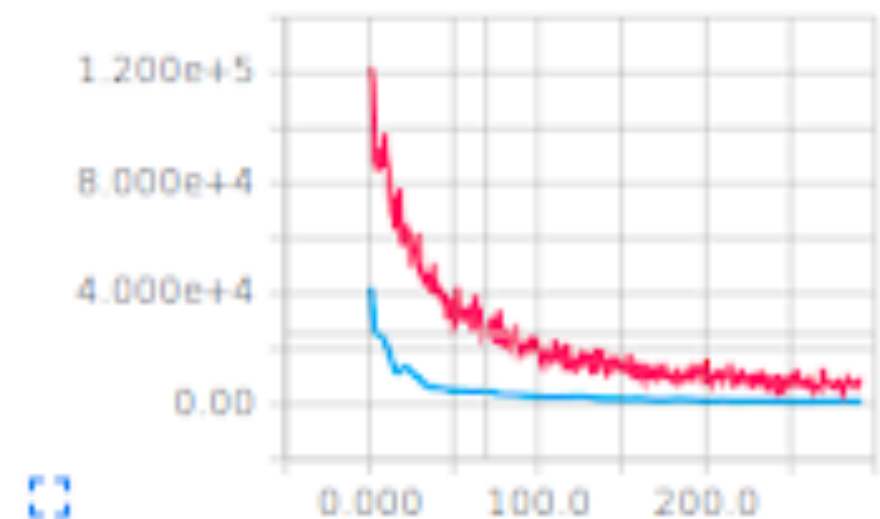
- Optimiser calculates gradients and updates once per batch
 - Batch size = number of samples
 - impacts memory requirements
 - impacts execution speed
- Repeat batches for whole training set
 - One epoch = all training samples
- Repeat epochs until convergence

accuracy



training —
validation —


loss



Loss and Activation Functions: Regression

Loss function options:

- `mean_squared_error`
- `mean_absolute_error`




Activation function for final node:

- `none` (unconstrained)
- `ReLU` (positive range)
- `sigmoid` (limited range)

Loss and Activation Functions: Classification

Loss function options:

- binary_crossentropy
- categorical_crossentropy
- sparse_categorical_crossentropy



Each prefers extreme values (near 0 or 1)

Activation function for final node:

One-hot representation

$$\text{softmax: } p_j = \frac{\exp(y_j)}{\sum_k \exp(y_k)}$$

Summary

- Network architecture: fully connected layers, one to next
- Neurons: importance of nonlinear activation functions (ReLU)
- Parameters: weights and biases, number of parameters
- Loss functions, epochs, batches, optimisers
- Regression:
 - loss = mean squared error ; mean absolute error
 - activation = None, ReLU, sigmoid
- Classification:
 - loss = cross entropy variants
 - activation = softmax (with one-hot representation)