

Using Machine Learning Tools

Convolutional Neural Networks

University of Adelaide

Previously ...

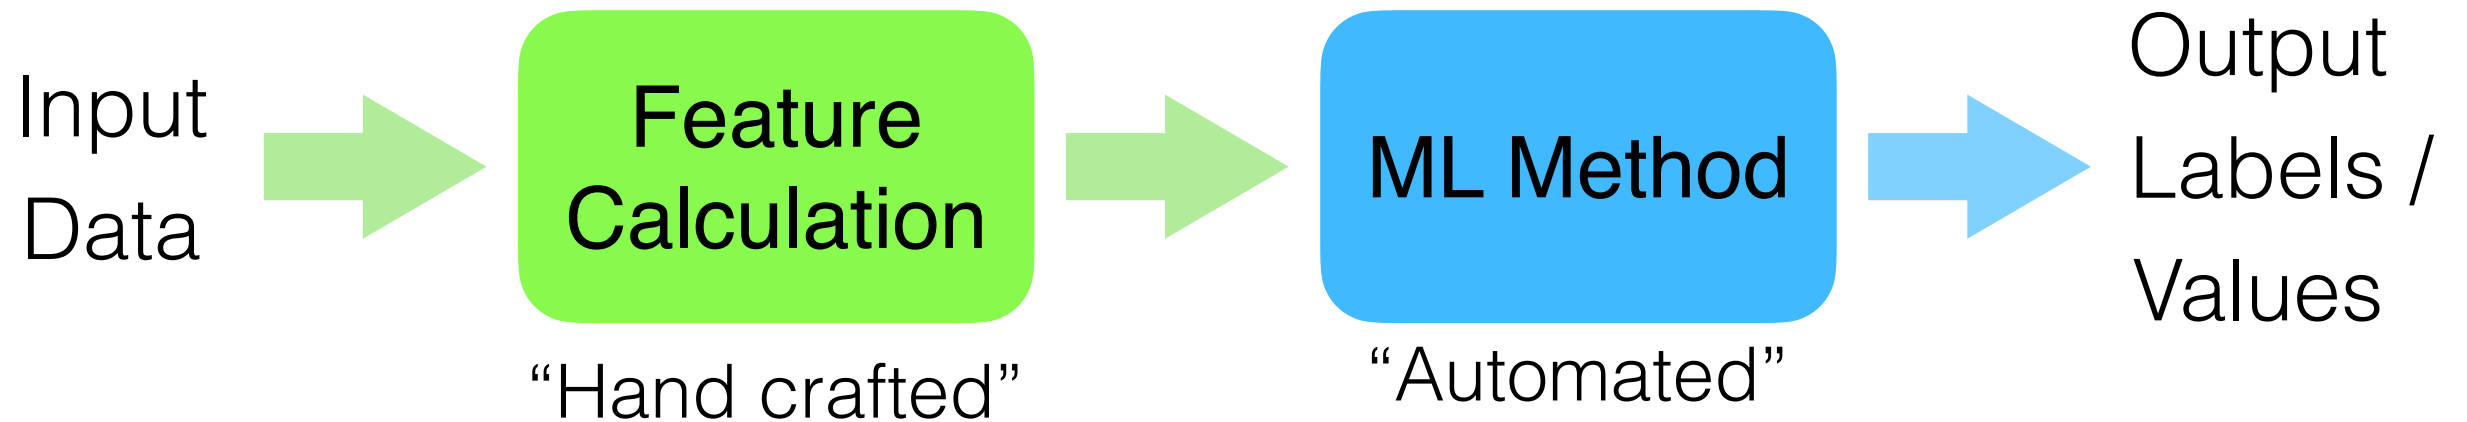
- Network architecture: fully connected layers, one to next
- Neurons: importance of nonlinear activation functions (ReLU)
- Parameters: weights and biases, number of parameters
- Loss functions, epochs, batches, optimisers
- Regression:
 - loss = mean squared error ; mean absolute error
 - activation = None, ReLU, sigmoid
- Classification:
 - loss = cross entropy variants
 - activation = softmax (with one-hot representation)

Today


- Convolutional Neural Networks (CNNs)
- Architecture:
 - convolution layer
 - pooling layer
 - stride
 - number of filters
- Convergence and overfitting
- Ensembles

Introduction to Neural Networks

Traditional Machine Learning



Deep Learning




Neural Network

Input Layer Hidden Layer 1 Hidden Layer 2 Output Layer

Input values as 1D vector

Fully Connected Layout

Each layer *only* connects to next layer



Bias inputs are not shown but implicit

Neural Network


Input
values as
1D vector

Reshaping into a 1D
vector means that
spatial relationships
are lost

Instead, use a linear
operation that works
with images and has
few parameters:


convolution

Input
Layer




Convolutional Neural Network

Input
Layer




Input values as
2D or 3D image

Hidden
Layer 1




Each layer *only*
connects to next layer

Hidden
Layer 2



...

Dense
Layers




Bias inputs
are not shown
but implicit

Output
Layer

Convolution Layer


Image



Filter

x derivative

-1	0	1
-2	0	2
-1	0	1



Convolution Layer

Image

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Convolution Layer

Image

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Convolution Layer

Image

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Convolution Layer

Image

	4	8	3	5	0
	2	4	5	3	0
	1	9	1	7	0
	7	6	0	0	0
	0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Convolution Layer

Image

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Convolution Layer

Image

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Convolution Layer

Image

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Convolution Layer

Image

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Convolution Layer

Image

	4	8	3	5	0
	2	4	5	3	0
	1	9	1	7	0
	7	6	0	0	0
	0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Convolution Layer

Image

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Convolution Layer

Image

4	8	3	5	0	
2	4	5	3	0	
1	9	1	7	0	
7	6	0	0	0	
0	0	0	0	0	

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Convolution Layer

Image

4	8	3	5	0	
2	4	5	3	0	
1	9	1	7	0	
7	6	0	0	0	
0	0	0	0	0	

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Activation is
also applied
to outputs


Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Bias inputs
are not shown
but implicit


Convolutional Neural Network

Input
Layer




Input values as
2D or 3D image

Hidden
Layer 1




Each layer *only*
connects to next layer

Hidden
Layer 2



...

Dense
Layers



Bias inputs
are not shown
but implicit

Output
Layer

Convolutional Neural Network


Input
Layer

Hidden
Layer 1

Have many
filters (kernels)

Each filter used
to make another
output image,
stacked together

Apply each
convolutional
filter in term



Input values as
2D or 3D image

Each layer *only*
connects to next layer

Convolutional Neural Network


Input
Layer

Hidden
Layer 1

Have many
filters (kernels)

Each filter used
to make another
output image,
stacked together

Apply each
convolutional
filter in term



Input values as
2D or 3D image

Each layer *only*
connects to next layer

Convolutional Neural Network


Input
Layer

Hidden
Layer 1

Have many
filters (kernels)

Each filter used
to make another
output image,
stacked together

Apply each
convolutional
filter in turn



e.g. depth = 4

Input values as
2D or 3D image

Each layer *only*
connects to next layer

Convolutional Neural Network

Input
Layer

Hidden
Layer 1


Hidden
Layer 2

Have many
filters (kernels)

Each filter used
to make another
output image,
stacked together

Apply each
convolutional
filter in term


Each filter acts
across all depths



Input values as
2D or 3D image

Each layer *only*
connects to next layer

Max Pooling Layer




Take maximum value within 2x2 filter

Move in steps of 2 (***stride = 2***)

Zero parameters

Max Pooling Layer




Take maximum value within 2x2 filter

Move in steps of 2 (***stride = 2***)


Zero parameters

Average Pooling is also an option

Convolutional Neural Network



Parameters



Number of parameters is $3 \times 3 = 9$

+ 1 bias input (implicit)


= 10 parameters

... for a scalar input image (depth=1)

NB: depth = 3 for colour channels

Independent of image size!

Convolutional Neural Network



Have many
filters (kernels)

Each filter used
to make another
output image,
stacked together

Apply each
convolutional
filter in term


For deeper stacks, the number of
parameters in *each filter* is:

depth * filter size + bias

e.g. $4 * 3 * 3 + 1 = 37$

Each filter acts
across all depths

Convolutional Neural Network

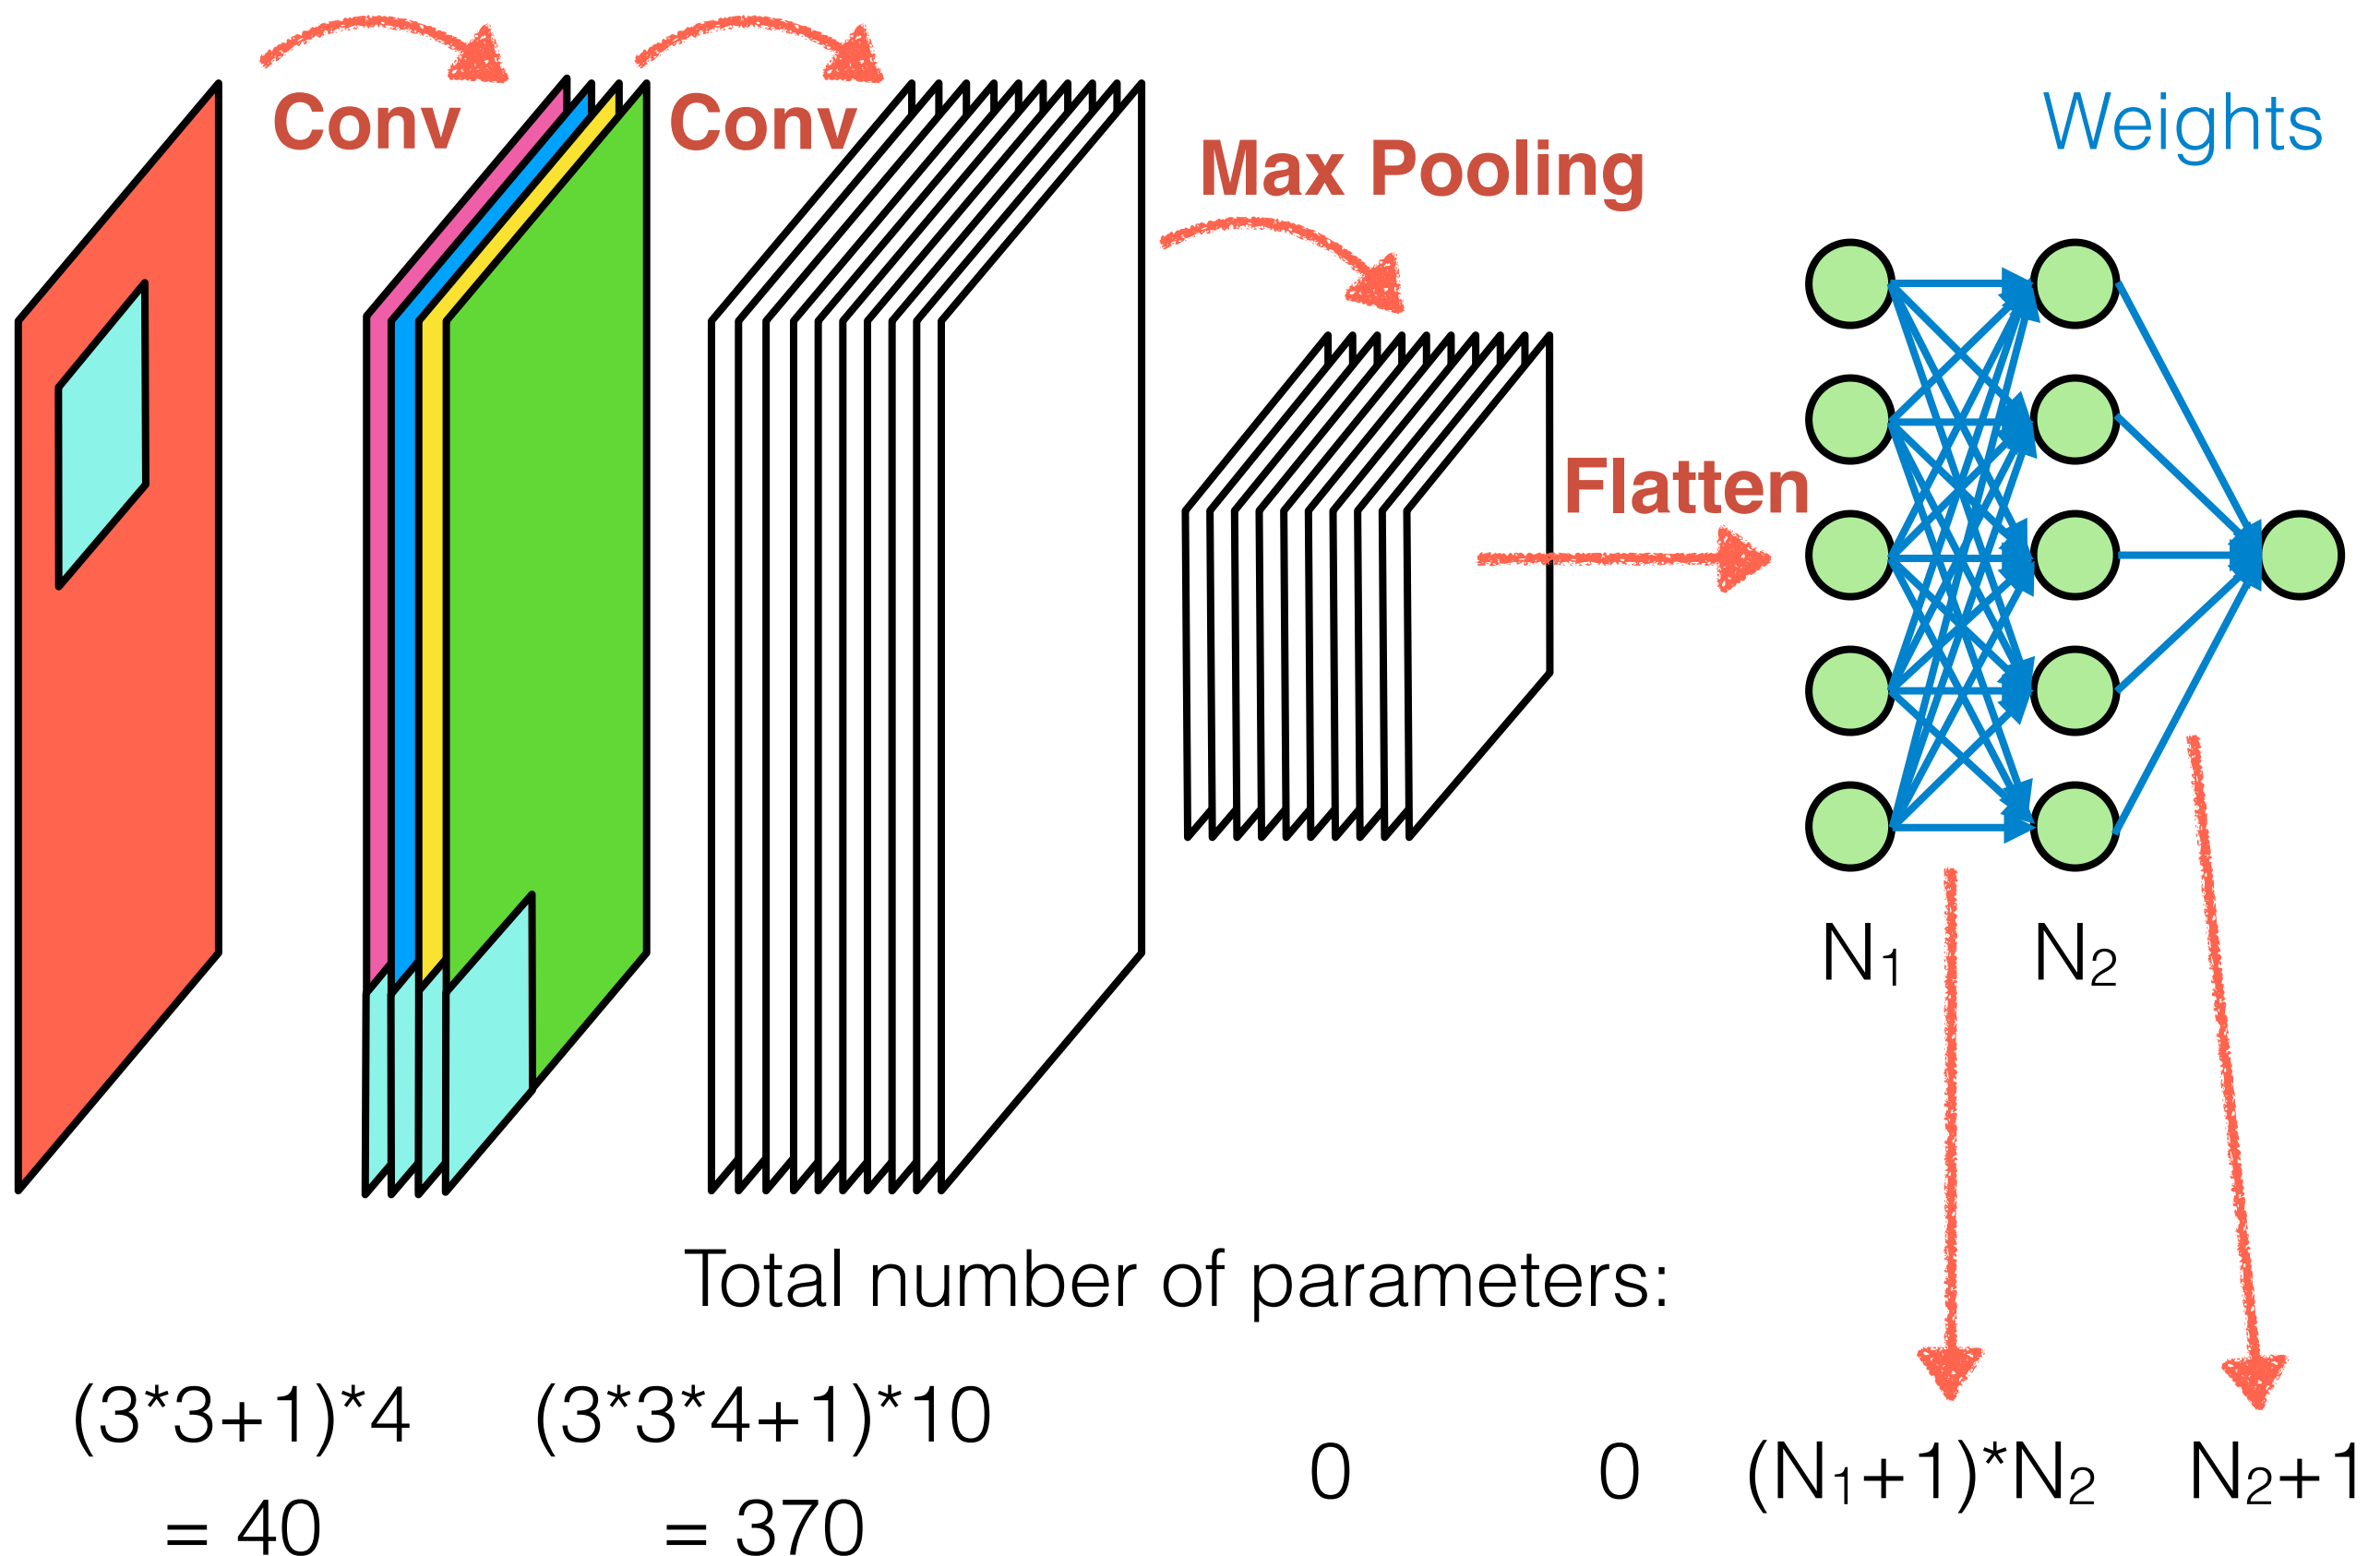


For deeper stacks, the number of
parameters in *each filter* is:


depth * filter size + bias

e.g. $4 * 3 * 3 + 1 = 37$

Parameters



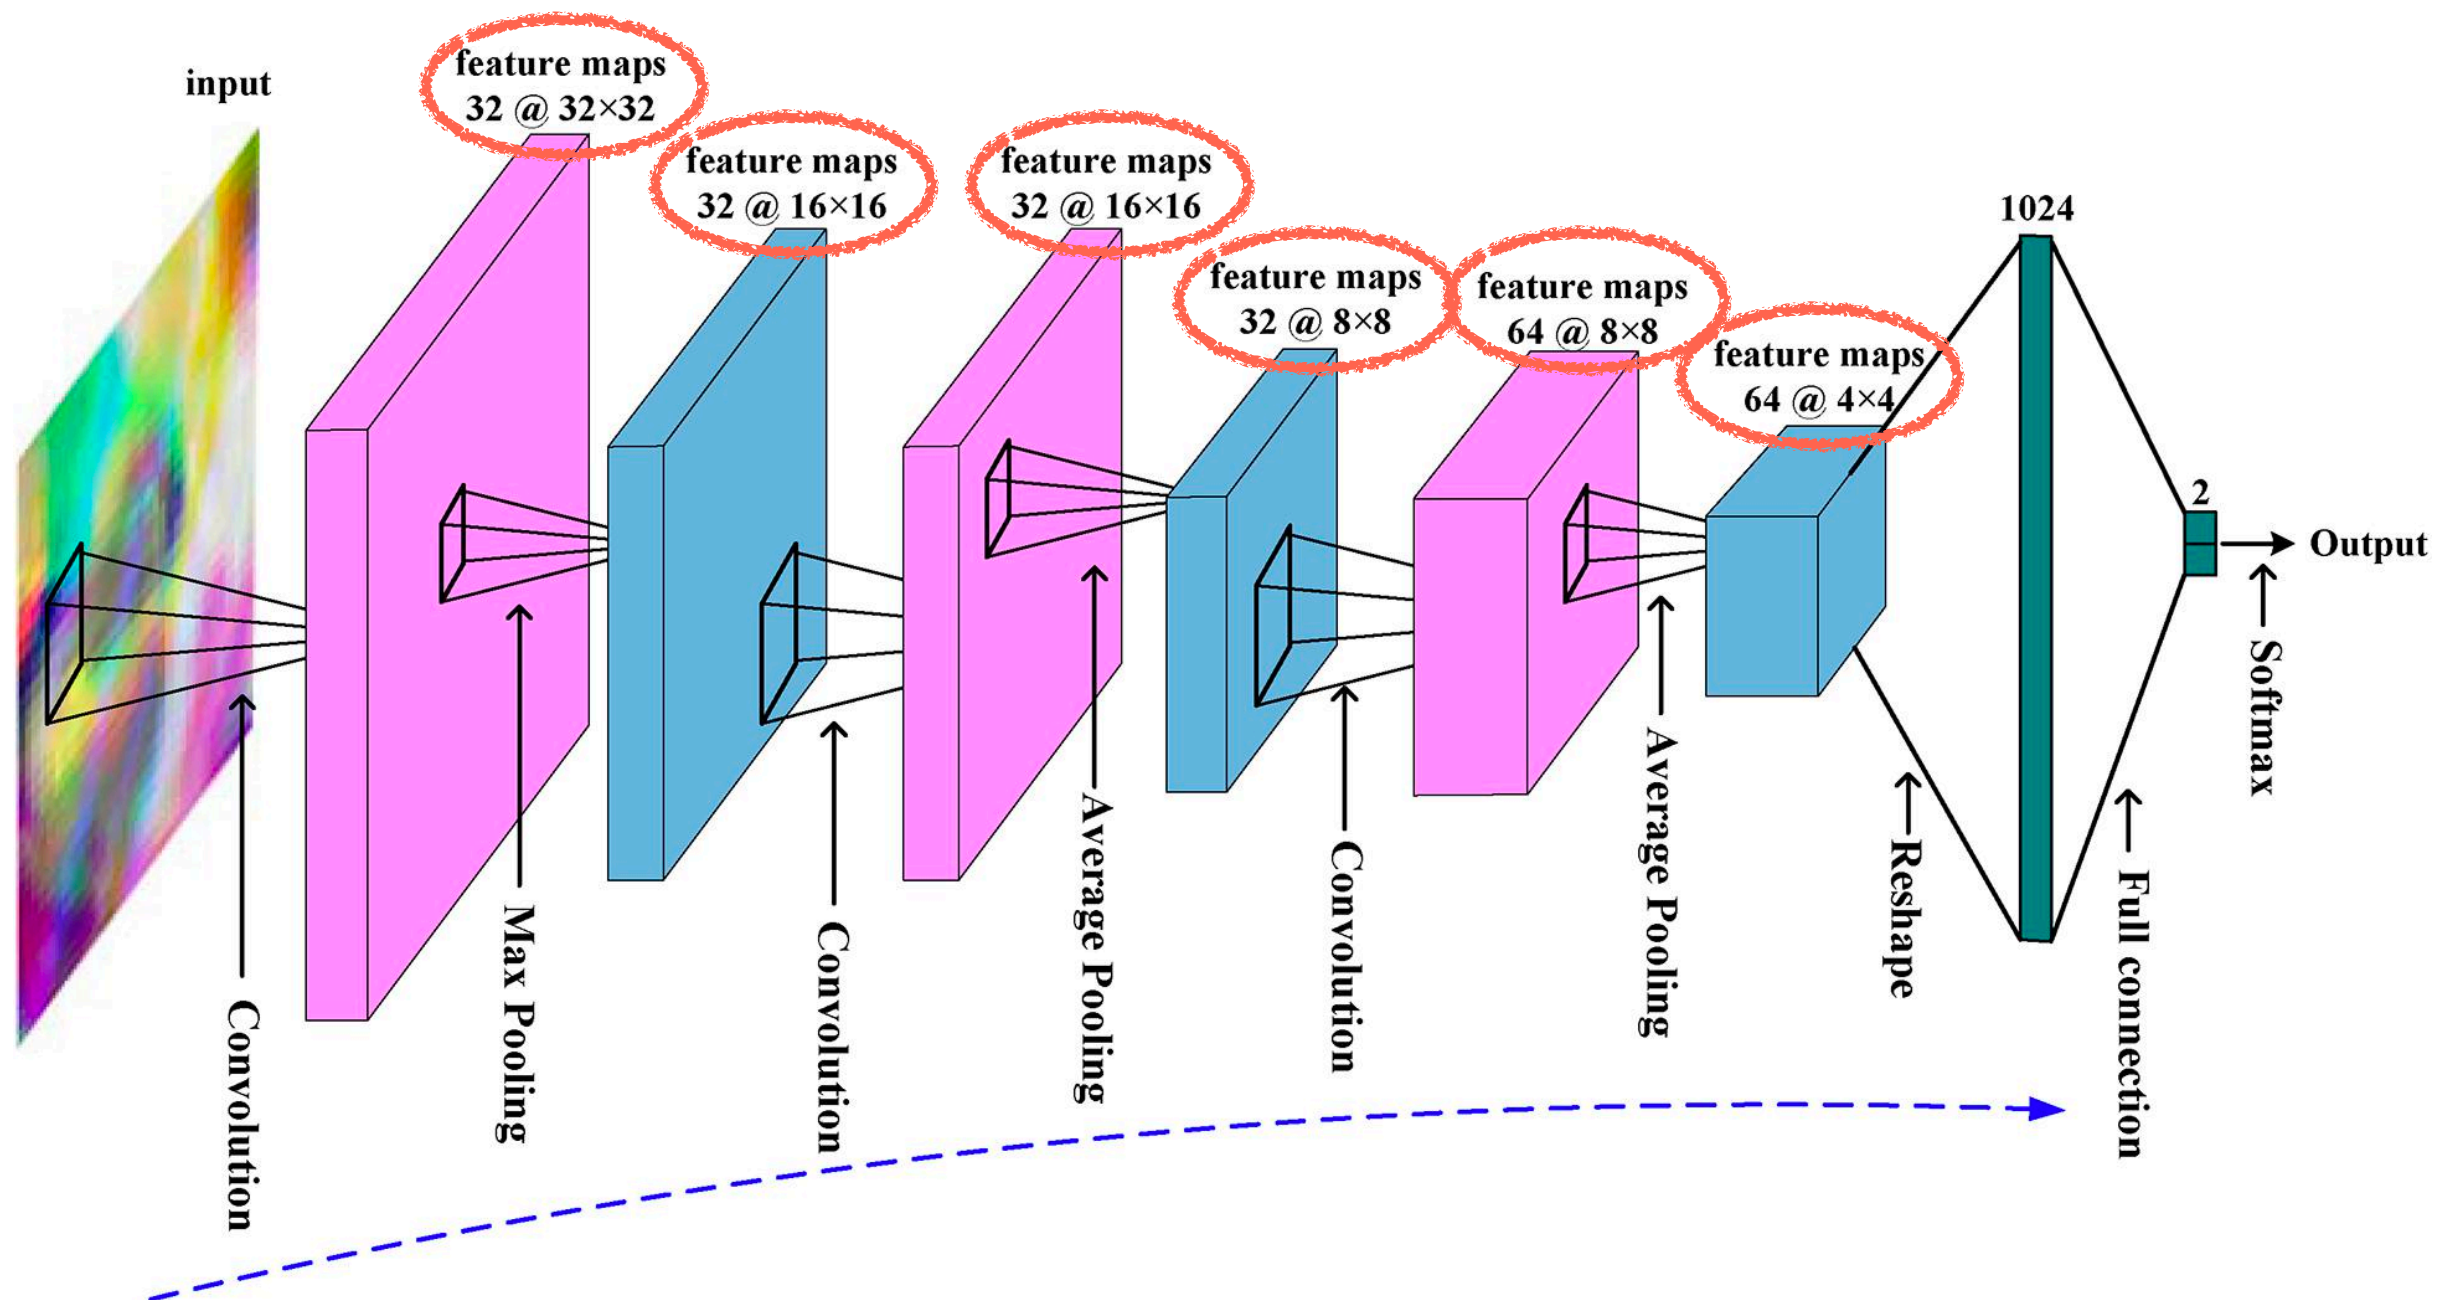
Example



Lin et al., *Convolutional Neural Networks-Based MRI Image Analysis for the Alzheimer's Disease Prediction From Mild Cognitive Impairment*, Front. Neurosci., 2018

Example

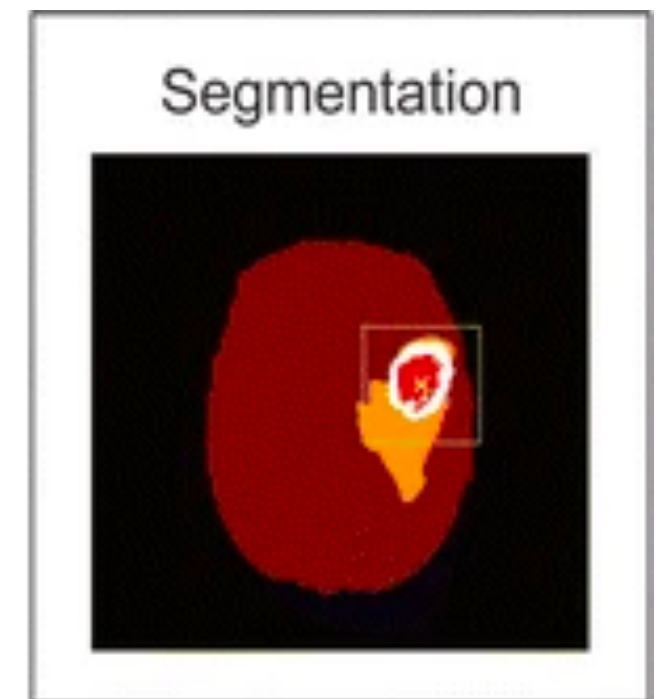
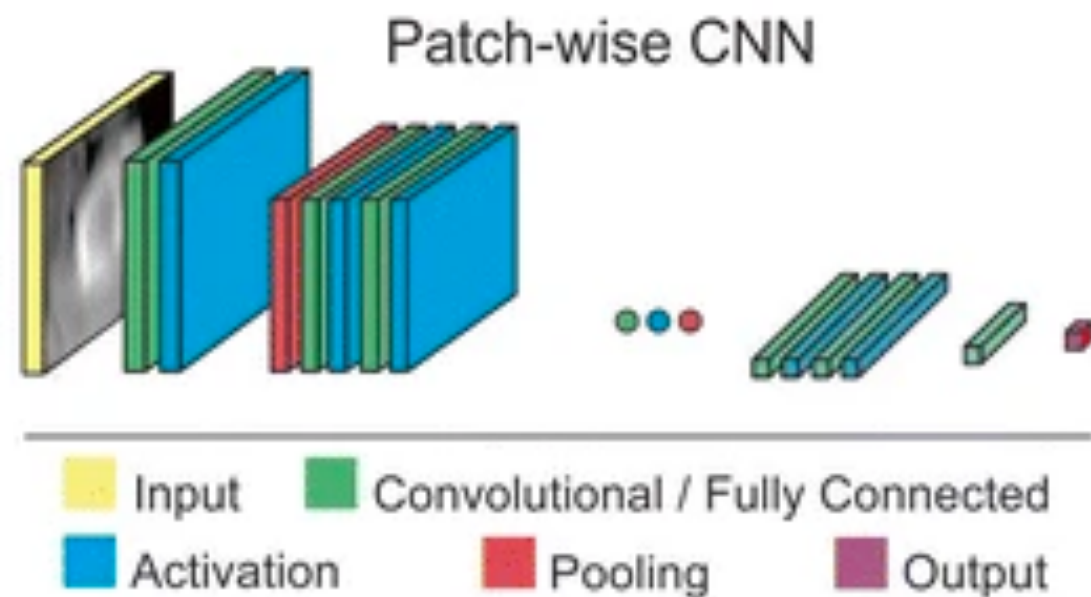

Lots of decisions about architecture



Lin et al., *Convolutional Neural Networks-Based MRI Image Analysis for the Alzheimer's Disease Prediction From Mild Cognitive Impairment*, Front. Neurosci., 2018

Example


Patch-based application for segmentation
Apply separately at each location



Convergence and Overfitting


- Batch size, epochs and convergence are the same as before
- Overfitting is more likely when number of parameters exceeds number of training data inputs
 - validation acc/loss much worse than training acc/loss
- Underfitting occurs when insufficient number of parameters or insufficient richness in data
 - training acc/loss is poor
- Can help guide network choices
- Empirical selection is still most common

accuracy



training —
validation —

loss



Ensembles

- Results of optimisation (training) will depend on many things (e.g. initialisation, number of epochs, order of data, etc.)
- Most of the time it will end up in a local minima
 - Though that might be good enough
- Can combine together different trained models and average outputs to get an improved model
 - Initialised differently
 - Trained differently
 - Different architectures / parameters
- Well studied in traditional machine learning
 - Random forests = ensemble of decision trees
 - This is one of the best models in traditional ML
- Approach is also common and powerful for Deep Learning

Summary

- Convolutional Neural Networks (CNNs)
- Architecture:
 - convolution layer (size & number of filters)
 - pooling layer (max or average - downsizes)
 - stride (can also apply to convolution)
 - number of filters and number of parameters
- Convergence and overfitting
- Ensembles

